


Enhancing Employee Attrition Prediction Using Hybrid Machine Learning Models

**Project report submitted in partial fulfillment of the requirements
for the award of the degree
BACHELOR OF TECHNOLOGY
in
COMPUTER SCIENCE& ENGINEERING (DATA SCIENCE)**

Submitted by

**P. Jacob Abraham Jane (Y22CD122)
K. Karthik (Y22CD071) N. Sai Sumanth Babu (Y22CD113)**

**Under the Guidance of
Mr. R.CH.S.N.P.Sairam
Assistant Professor**




R.V.R. & J.C. COLLEGE OF ENGINEERING (AUTONOMOUS)

Approved by AICTE- New Delhi, Accredited by NAAC A+ Grade

**Affiliated to Acharya Nagarjuna University,
Chowdavaram, Guntur - 19, Andhra Pradesh, India.**

MAY 2026

R.V.R. & J.C. COLLEGE OF ENGINEERING (AUTONOMOUS)
DEPARTMENT OF
COMPUTER SCIENCE & ENGINEERING (DATA SCIENCE)



CERTIFICATE

This is to certify that the project work titled “Enhancing Employee Attrition Prediction using Hybrid Machine Learning Models” is the work done by **P. Jacob Abraham Jane (Y22CD122)**, **K. Karthik (Y22CD071)** and **N. Sai Sumanth Babu (Y22CD113)** under my supervision and submitted in partial fulfillment of the requirements for the award of the degree, B.Tech. in Computer Science and Engineering (Data Science), during the Academic Year 2025-2026.

Mr. R.CH.S.N.P.Sairam
Assistant Professor,
Project Guide

Dr. M.V.P.Chandra Sekhara Rao
Professor & Head

External Examiner

DECLARATION

We **P. Jacob Abraham Jane (Y22CD122)** and **K. Karthik (Y22CD071)** and **N .Sai Sumanth Babu (Y22CD113)** hereby declare that the project report titled **“Enhancing Employee Attrition Prediction Using Hybrid Machine Learning Models”** under the guidance of **Mr.R.CH.S.N.P.Sairam.** is submitted in partial fulfillment of the requirements for the award of the degree of Bachelor of Technology in Computer Science and Engineering (Data Science). This is a record of bonafide work carried out by us and the results embodied in this project have not been reproduced or copied from any source. The results embodied in this project have not been submitted to any other university for the award of any other degree.

P. Jacob Abraham Jane (Y22CD122)

K. Karthik (Y22CD071)

N. Sai Sumanth Babu (Y22CD113)

Place: Guntur

Date:

ACKNOWLEDGMENT

The successful completion of any task would be incomplete without proper suggestion, guidance, and environment. Combination of these three factors acts like a backbone to our project work “**Enhancing Employee Attrition Prediction Using Hybrid Machine Learning Models**”.

We are profoundly pleased to express our deep sense of gratitude and respect towards the management of the **R.V.R. & J.C. College of Engineering** for providing the resources to complete the project.

We are very much thankful to **Dr. Kolla Srinivas**, Principal, **R.V.R. & J.C. College of Engineering** for allowing us to deliver the project successfully.

We are greatly indebted to **Dr. M.V.P.Chandra Sekhara Rao**, Professor, & Head of the department, Computer Science and Engineering (Data Science), **R.V.R. & J.C. College of Engineering**, providing the laboratory facilities fully as and when required and for giving us the opportunity to carry the project work in the college.

We are also thankful to our Project Coordinator **Dr. Popuri Srinivasa Rao**, Associate Professor who helped us in each step of our Project.

We are very much thankful to our Project guide **Mr. R.CH.S.N.P.Sairam**, Assistant Professor who has/her inspired us to take this project and gave us valuable guidance to complete the project successfully.

We would be thankful to **all the staff** of the Department of Computer Science and Engineering (Data Science) for their valuable suggestions, guidance and constructive ideas in each and every step, which was indeed of great help towards the successful completion of our project.

P. Jacob Abraham Jane (Y22CD122)

K. Karthik (Y22CD071)

N. Sai Sumanth Babu (Y22CD113)

ABSTRACT

Employee attrition is a critical challenge for organizations, leading to increased operational costs, reduced productivity, and loss of skilled workforce. Traditional prediction approaches often rely on single machine learning models, which struggle to handle class imbalance, complex feature interactions, and lack interpretability required for effective decision-making. To address these limitations, this project proposes a hybrid machine learning framework for employee attrition prediction that integrates advanced preprocessing, feature optimization, ensemble learning, and explainable artificial intelligence techniques. The proposed system utilizes the Synthetic Minority Oversampling Technique (SMOTE) to handle class imbalance and improve minority-class representation. A Genetic Algorithm (GA) is employed for feature selection to identify the most relevant attributes and reduce redundancy in high-dimensional HR datasets. The optimized features are then used to train an ensemble model combining LightGBM and CatBoost classifiers using a soft-voting strategy, enabling the model to capture complex nonlinear relationships and improve predictive robustness. To enhance transparency and interpretability, SHAP (SHapley Additive exPlanations) is incorporated to provide feature-level insights and explain individual predictions. This allows the system to identify key factors influencing employee attrition and supports data-driven human resource decision-making. The proposed framework provides a scalable, efficient, and interpretable solution for employee attrition prediction, making it suitable for real-world HR analytics applications. It enables organizations to proactively identify at-risk employees and implement effective retention strategies, thereby improving workforce stability and organizational performance..

TABLE OF CONTENTS

TITLE	PAGENO.
ABSTRACT	I
TABLE OF CONTENTS	II
LIST OF FIGURES	IV
LIST OF TABLES	V
CHAPTER 1: INTRODUCTION	
1.1. Introduction	2
1.2. Problem Statement	2
1.3. Objectives of the Study	3
CHAPTER 2: LITERATURE SURVEY	
2.1. Literature Survey	5
CHAPTER 3: SYSTEM ANALYSIS & FEASIBILITY STUDY	
3.1. Existing System	9
3.1.1. Limitations of Existing System	10
3.2. Proposed System	10
3.2.1. Advantages of Proposed System	11
3.2.2. Dataset	11
3.2.3. Data Preprocessing	12
3.2.3.1. Data Cleaning	13
3.2.3.2. Normalization	13
3.2.3.3. Categorical Encoding	14
3.2.3.4. Feature Scaling	14
3.2.3.5. Stratified Train-Test Split	14
3.2.3.6. Smote Oversampling	15
3.3. Model Building	15
3.3.1. Genetic Algorithm-Based Feature Selection	17
3.3.2. Ensemble Classification Model	17
3.3.3. Prediction Layer	18
3.3.4. Performance Evaluation	18
3.3.5. Explainability using SHAP	19
3.3.6. Loss Function: Binary Cross-Entropy	19
3.4. Model Validation	20
3.5. Evaluation Metrics	21
3.6. Feasibility Study	
3.6.1. Economic Feasibility	22
3.6.2. Operational Feasibility	22
3.6.3. Technical Feasibility	23

CHAPTER 4: SYSTEM REQUIREMENTS	
4.1. Functional Requirements	25
4.2. Non-Functional Requirements	26
4.2.1. Software Requirements	26
4.2.2. Hardware Requirements	27
CHAPTER 5: DESIGN	
5.1. System Design	29
5.2. UML Diagrams	
5.2.1. Class Diagram	33
5.2.2. Sequence Diagram	34
5.2.3. Use Case Diagram	35
5.2.4. Workflow Diagram	36
CHAPTER 6: IMPLEMENTATION	
6.1. Implementation	38
CHAPTER 7: RESULTS	
7.1. Model Results	46
7.2. Results Comparison	47
7.3. ROC Curve Analysis	49
7.4. Confusion Matrix Analysis	49
7.5. Feature Importance Analysis	50
7.6. Model Explainability using SHAP	51
CHAPTER 8: SOCIAL IMPACT	
8.1. Social Impact	54
CHAPTER 9: CONCLUSION & FUTURE WORK	
9.1. Conclusion & Future Work	57
CHAPTER 10: BIBLIOGRAPHY	60

LIST OF FIGURES

FigureNo.	Figure Description	PageNo.
3.1	Proposed Architecture	16
5.1	Block Diagram of Proposed Hybrid Employee Attrition Prediction	29
5.2.1	UML Class Diagram of Employee Attrition Prediction System	33
5.2.2	UML Sequence Diagram for System Interaction Flow	34
5.2.3	UML Use Case Diagram of Employee Attrition Prediction System	35
5.2.4	Workflow Diagram of Proposed Hybrid Machine Learning Model	36
7.1	ROC Curve Comparison	42
7.2	Comparison of Proposed and Existing Models	44
7.3	Comparative Performance of All Machine Learning Models	44
7.4	Confusion Matrix	46
7.5	Feature Importance	47
7.6	SHAP Feature Importance	48

LIST OF TABLES

TableNo.	Table Description	PageNo.
2.1	Survey on Recent Issues	6
7.1	Comparison of Classification Metrics for Existing and Proposed Models	43

Chapter 1

Introduction

1. INTRODUCTION

1.1 Introduction

Employee attrition is a major challenge for organizations, affecting productivity, operational stability, and overall performance. High turnover leads to increased recruitment costs, loss of skilled employees, and workflow disruptions, making early prediction of attrition important for effective Human Resource (HR) management [11].

With the rise of data-driven approaches, machine learning has become essential for analyzing employee data and predicting attrition. HR datasets contain diverse attributes such as demographics, job roles, compensation, and satisfaction metrics, creating complex relationships. Traditional methods often fail to capture these patterns, whereas machine learning models, particularly ensemble and boosting techniques such as LightGBM and XGBoost, can effectively model such complexities [8][10].

However, challenges such as class imbalance and lack of interpretability still exist. Employees leaving the organization form a minority class, leading to biased predictions, while many models act as black-box systems, limiting trust in HR decisions [15][5]. To overcome these issues, this project proposes a hybrid framework integrating SMOTE, Genetic Algorithms (GA), LightGBM, CatBoost, and SHAP to improve prediction reliability and provide interpretable insights [6][9].

1.2 Problem Statement

Despite advancements in machine learning, existing employee attrition prediction systems face several limitations in real-world HR datasets. Traditional single-model approaches often struggle to capture complex relationships in high-dimensional data, leading to poor generalization across diverse workforce patterns [12]. A major issue is class imbalance, where employees leaving the organization form a minority class, causing models to bias toward majority retention cases and reducing the effectiveness

of attrition detection [15]. Additionally, the lack of proper feature selection results in redundant attributes, increasing model complexity and reducing efficiency [14].

Another key limitation is the lack of interpretability in many machine learning models. High-performing models often act as black-box systems, making it difficult for HR professionals to understand and trust predictions [5]. Moreover, relying on a single algorithm limits predictive robustness, as different models capture different patterns within the data. These limitations reduce the practical usability of existing systems in real-world HR decision-making environments.

Furthermore, many existing approaches do not provide a unified solution that addresses imbalance handling, feature optimization, prediction accuracy, and interpretability. As a result, organizations lack reliable tools to effectively identify employees at risk of attrition. Therefore, there is a need for an integrated approach that improves prediction reliability through ensemble learning while providing clear, interpretable insights for informed decision-making [8][6].

1.3 Objectives of the Study

- To develop an efficient and interpretable hybrid machine learning model for employee attrition prediction.
- To implement a robust preprocessing pipeline including handling missing values, encoding categorical data, and normalizing features [12].
- To address class imbalance using SMOTE and optimize features using Genetic Algorithm (GA) for improved model performance [15][18].
- To build a hybrid ensemble model using LightGBM and CatBoost with a soft-voting mechanism for enhanced accuracy and robustness [8][6].
- To incorporate SHAP-based explainability for interpreting predictions and enabling data-driven HR decision-making [9][11].
- To improve overall prediction reliability and support proactive employee retention strategies in organizations..

Chapter 2

Literature Survey

2. LITERATURE SURVEY

2.1 Literature Survey

Employee attrition prediction has become an important research area in Human Resource (HR) analytics, with the objective of improving workforce stability and organizational performance. Early studies primarily focused on traditional machine learning techniques such as Logistic Regression, Support Vector Machines (SVM), and Random Forest classifiers. These models utilized employee demographic, job-related, and behavioral data to predict turnover. Although they provided baseline performance, they often struggled to capture complex nonlinear relationships in high-dimensional HR datasets, limiting their effectiveness in real-world scenarios [17].

To overcome these limitations, researchers introduced ensemble learning and boosting algorithms such as XGBoost and LightGBM. Jain and Nayyar (2018) demonstrated that boosting methods significantly outperform traditional classifiers by effectively modeling feature interactions in structured data [7]. Similarly, recent hybrid approaches combining Genetic Algorithm (GA) with LightGBM have shown improved predictive performance by optimizing feature selection and enhancing model robustness [1]. However, these models often rely on single classifiers and lack ensemble diversity and interpretability mechanisms.

Recent research has also explored deep learning approaches for employee attrition prediction. Ahmed (2023) utilized neural network-based models to capture complex patterns within employee data and achieved improved prediction accuracy [2]. Despite these advantages, such models often lack interpretability and act as black-box systems, making them less suitable for HR decision-making environments where transparency is critical [5]. Additionally, deep learning models require higher computational resources and may not generalize well across different datasets.

Another key research direction focuses on addressing class imbalance in attrition datasets. Techniques such as the Synthetic Minority Oversampling Technique (SMOTE) have been widely adopted to improve minority-class representation and

enhance recall performance [15]. However, many existing studies apply imbalance handling, feature selection, and classification independently rather than integrating them into a unified framework, which limits overall model efficiency and robustness. The proposed work (2026) develops an integrated hybrid framework combining SMOTE for imbalance handling, GA for feature selection, and a LightGBM–CatBoost ensemble using soft voting [8][6]. It also incorporates SHAP for explainability [9]. This approach improves prediction accuracy, interpretability, and practical applicability for employee attrition prediction.

Table 2.1: Survey on Recent Issues

Authors (Year)	Title of Paper	Dataset	Techniques Used
Sirawit Saiwaeo et al. (2023) [1]	Developing a Hybrid Machine Learning Model for Employee Turnover Prediction Using LightGBM + GA	HR Analytics & Organizational datasets	GA-based feature selection with LightGBM improves AUC and F1-score.
Ahmed (2023) [2]	Predicting Employee Attrition Using Artificial Neural Networks	Organizational HR dataset	Neural networks capture non-linear patterns and improve prediction accuracy.
J. Judrus et al. (2021) [3]	Machine Learning-Based Solution for Predicting Voluntary Employee Turnover	Organizational dataset	Compared multiple ML classifiers for attrition prediction.
Z. Zhao et al. (2019) [4]	Maximum Relevance and Minimum Redundancy Feature Selection Methods	General datasets	Proposed feature selection method to improve model efficiency.
Z. C. Lipton (2018) [5]	The Mythos of Model Interpretability	Conceptual study	Highlights challenges in interpreting black-box ML models.
L.Prokhorenkova et al. (2018) [6]	CatBoost: Unbiased Boosting with Categorical Features	Benchmark datasets	Introduced CatBoost for handling categorical features effectively.

Jain & Nayyar (2018) [7]	Predicting Employee Attrition using XGBoost	HR Dataset	Demonstrated effectiveness of XGBoost for attrition prediction.
G. Ke et al. (2017) [8]	LightGBM: A Highly Efficient Gradient Boosting Decision Tree	Benchmark datasets	Proposed LightGBM for fast and efficient gradient boosting.
Lundberg & Lee (2017) [9]	A Unified Approach to Interpreting Model Predictions	General datasets	Introduced SHAP for model explainability.
Chen & Guestrin (2016) [10]	XGBoost: A Scalable Tree Boosting System	Large-scale datasets	Developed scalable boosting algorithm for high performance.
Davenport et al. (2010) [11]	Competing on Talent Analytics	HR analytics	Emphasized importance of data-driven HR decision-making.
Hastie et al. (2009) [12]	The Elements of Statistical Learning	General datasets	Foundation for statistical and ML techniques.
Fawcett (2006) [13]	An Introduction to ROC Analysis	Classification datasets	Explained ROC-AUC for model evaluation.
Guyon & Elisseeff (2003) [14]	An Introduction to Variable and Feature Selection	General datasets	Discussed importance of feature selection techniques.
Chawla et al. (2002) [15]	SMOTE: Synthetic Minority Over-sampling Technique	Imbalanced datasets	Proposed SMOTE for handling class imbalance.
Friedman (2001) [16]	Gradient Boosting Machine	General datasets	Introduced boosting techniques for predictive modeling.
Breiman (2001) [17]	Random Forests	General datasets	Proposed ensemble learning using decision trees.
Goldberg (1989) [18]	Genetic Algorithms in Search, Optimization and Machine Learning	General optimization problems	Introduced GA for optimization and feature selection.

Chapter 3

System Analysis & Feasibility

Study

3. SYSTEM ANALYSIS & FEASIBILITY STUDY

3.1 Existing System

Traditionally, employee attrition prediction systems rely on basic statistical methods and conventional machine learning algorithms such as Logistic Regression, Decision Trees, and Support Vector Machines (SVM). These models utilize employee-related data such as demographics, job roles, and performance metrics to predict whether an employee is likely to leave the organization. While these approaches are simple and easy to implement, they often struggle to capture complex nonlinear relationships present in high-dimensional HR datasets. As a result, their predictive performance and generalization ability remain limited in real-world scenarios [12].

In recent years, advanced machine learning techniques such as Random Forest, XGBoost, and LightGBM have been introduced to improve prediction accuracy. These models are capable of handling structured data more effectively and capturing intricate feature interactions [8][10]. However, many existing systems still rely on single-model architectures and do not integrate multiple learning approaches. Additionally, most of these models do not adequately address class imbalance, where employees leaving the organization form a minority class. This leads to biased predictions toward the majority class and reduces the effectiveness of attrition detection [15].

Furthermore, many existing systems lack proper feature selection mechanisms, resulting in the inclusion of redundant or irrelevant attributes that increase model complexity and reduce efficiency [14]. Another major limitation is the lack of interpretability, as many high-performing models operate as black-box systems [5]. This makes it difficult for HR professionals to understand the reasoning behind predictions and limits their adoption in decision-making processes. Consequently, existing attrition prediction systems often produce less reliable, less interpretable, and less practical outcomes for real-world HR analytics applications. These limitations highlight the need for integrated and explainable solutions to improve both accuracy and usability in practice.

3.1.1 Limitations of existing system

Traditional employee attrition prediction systems, particularly those based on single machine learning models, face several limitations when applied to real-world HR datasets. These systems often rely on basic algorithms such as Logistic Regression, Decision Trees, or standalone boosting models, which struggle to capture complex nonlinear relationships in high-dimensional data [12]. One of the major challenges is class imbalance, where employees leaving the organization form a minority class, causing models to bias toward majority retention cases and resulting in poor detection of actual attrition instances [15].

Additionally, many existing systems lack proper feature selection mechanisms and include redundant or irrelevant attributes, which increases model complexity and reduces efficiency [14]. Most models also do not integrate multiple learning techniques, limiting predictive robustness. Another significant limitation is the lack of interpretability, as many high-performing models act as black-box systems, making it difficult for HR professionals to understand and trust the predictions [5]. These limitations reduce the effectiveness and practical usability of existing attrition prediction systems..

3.2 Proposed System

The proposed system presents an advanced and integrated approach for employee attrition prediction by combining multiple machine learning techniques into a unified framework. The system is designed to improve prediction accuracy and reliability by incorporating data preprocessing, class imbalance handling, feature optimization, ensemble learning, and explainability. A structured pipeline is implemented to process HR data, ensuring consistency and compatibility for model training [12].

The proposed approach utilizes the Synthetic Minority Oversampling Technique (SMOTE) to address class imbalance by improving minority-class representation [15]. Genetic Algorithm (GA) is employed for feature selection to identify the most relevant attributes and reduce redundancy in the dataset [18]. A hybrid ensemble model is developed by combining LightGBM and CatBoost classifiers using a soft-

voting mechanism, allowing the system to capture complex feature interactions and improve predictive robustness [8][6]. Furthermore, SHAP-based explainability is incorporated to provide transparent insights into model predictions, enabling HR professionals to understand the key factors influencing employee attrition [9]. This integrated approach overcomes the limitations of existing systems and provides a more accurate, reliable, and interpretable solution for real-world HR analytics applications.

3.2.1 Advantages of proposed system:

The proposed system employs a hybrid machine learning framework that integrates data preprocessing, class imbalance handling, feature optimization, ensemble learning, and explainability, enabling more accurate and reliable employee attrition prediction. The use of SMOTE effectively addresses class imbalance by improving minority-class representation, which enhances the model's ability to detect actual attrition cases [15]. Genetic Algorithm (GA)-based feature selection helps in identifying the most relevant attributes while eliminating redundant features [18].

Additionally, the combination of LightGBM and CatBoost classifiers in a soft-voting ensemble enhances predictive robustness by leveraging the strengths of multiple learning algorithms [8][6]. This approach enables the model to capture complex nonlinear relationships within HR datasets more effectively. The incorporation of SHAP-based explainability provides clear insights into feature contributions, allowing HR professionals to understand and trust the predictions [9]. Overall, the proposed system improves prediction accuracy, enhances interpretability, and offers a scalable and practical solution for real-world HR analytics applications.

3.2.2 Dataset

The dataset used for this study is the publicly available IBM HR Analytics dataset, which is widely utilized in employee attrition research [3]. The dataset consists of structured employee records containing demographic information, job-related attributes, compensation details, and performance indicators. Each record represents

an individual employee profile, with the target variable indicating whether the employee has left the organization or not.

The dataset includes diverse features such as age, education, job role, job satisfaction, work-life balance, monthly income, and years of experience, enabling the model to capture multiple factors influencing employee turnover. Since the dataset exhibits class imbalance, with fewer instances of employee attrition compared to retention, it provides a suitable scenario for evaluating imbalance-handling techniques such as SMOTE [15]. The data is preprocessed to remove inconsistencies, handle missing values, and transform categorical variables into numerical formats, ensuring high-quality input for model training and evaluation.

3.2.3 Data Pre-Processing

Data preprocessing is a crucial step in any machine learning-based prediction system, as it ensures that the input data is clean, consistent, and suitable for model training. In the proposed employee attrition prediction system, the HR dataset undergoes multiple preprocessing steps to improve data quality, reduce noise, and enhance model performance. These steps help in preparing structured data for effective learning and ensure better generalization in real-world scenarios [12].

The preprocessing phase includes handling missing values, removing duplicate records, and transforming raw data into a structured format. Categorical variables such as job role, department, and education are encoded into numerical representations to make them suitable for machine learning algorithms. Additionally, numerical features are normalized using techniques such as Min-Max scaling to ensure all features lie within a uniform range, which helps improve model convergence and stability during training [12].

Furthermore, irrelevant or redundant features are identified and eliminated to reduce dimensionality and improve computational efficiency. Proper preprocessing also minimizes noise in the dataset, allowing the model to learn meaningful patterns more effectively. Since the dataset exhibits class imbalance, preprocessing prepares the data for imbalance handling using SMOTE, which enhances the representation of

minority attrition cases and improves predictive performance [15]. Overall, these preprocessing steps play a vital role in building a robust, accurate, and scalable employee attrition prediction system.

3.2.3.1 Data Cleaning

The raw HR dataset may contain missing values, duplicate records, and irrelevant attributes that can negatively impact model performance. Therefore, preprocessing begins with removing duplicate entries and handling missing values using appropriate techniques such as mean imputation for numerical features. This ensures data consistency and prevents bias during model training [12].

In addition, irrelevant or non-informative columns are removed to retain only meaningful attributes that contribute to prediction. Cleaning the dataset also helps in reducing noise and improving the overall quality of input data. Proper data cleaning plays a vital role in enhancing model accuracy and efficiency by ensuring that the learning algorithms are trained on reliable and well-structured data [14].

3.2.3.2 Normalization

Numerical features such as age, income, and years of experience are normalized to a fixed range to stabilize model training and improve convergence. Min-Max normalization is applied to scale values between 0 and 1, ensuring uniformity across all numerical inputs. This prevents features with larger ranges from dominating the learning process and improves the overall performance of machine learning models.

By bringing all numerical features to a common scale, normalization enhances model stability, speeds up the learning process, and ensures that each feature contributes equally during training. This step is particularly important for algorithms that are sensitive to feature magnitude, as it helps in achieving better accuracy and efficient optimization [12]. It also reduces the chances of numerical instability during computation and supports faster convergence of gradient-based learning algorithms. Additionally, normalization improves consistency across different datasets, making the model more robust and reliable..

3.2.3.3 Categorical Encoding

Since HR datasets contain categorical attributes such as job role, department, and education level, these features are converted into numerical format using label encoding. This transformation allows machine learning models to process categorical data effectively while preserving the relationships between different categories [12].

Categorical encoding is essential because most machine learning algorithms require numerical input for computation. By converting categorical values into numerical labels, the dataset becomes suitable for model training without losing important information. This step ensures efficient processing of data and contributes to improved model performance and accuracy [12].

3.2.3.4 Feature Scaling

Feature scaling is performed to ensure that all input features contribute equally to the model. Min-Max scaling is applied across all numerical attributes to bring them to a comparable range. This step improves model stability, enhances convergence speed, and ensures that no single feature disproportionately influences the prediction process [12].

By standardizing the range of features, feature scaling helps machine learning algorithms learn more effectively and reduces bias caused by varying feature magnitudes. It also improves the efficiency of optimization techniques, leading to better accuracy and overall model performance [12].

3.2.3.5 Stratified Train-Test Split

The dataset is divided into training and testing sets using a stratified train-test split approach. This technique ensures that the proportion of attrition and non-attrition cases is maintained consistently in both training and testing datasets. Maintaining class distribution is particularly important in imbalanced datasets, where the minority class (employees leaving) needs to be properly represented [15].

By preserving the original class distribution, stratified sampling improves the reliability and fairness of model evaluation. It helps the model learn patterns effectively from the training data while ensuring that the testing phase reflects real-world data scenarios. This results in more accurate performance estimation and better generalization of the model [15].

3.2.3.6 Smote Oversampling

To address class imbalance in the dataset, the Synthetic Minority Oversampling Technique (SMOTE) is applied to the training data. In employee attrition datasets, the number of employees leaving the organization is typically much lower than those who stay, which can bias the model toward the majority class. SMOTE overcomes this issue by generating synthetic samples for the minority class through interpolation between existing minority instances [15].

This technique helps in balancing the dataset without simply duplicating existing records, thereby reducing overfitting and improving the model's ability to learn meaningful patterns related to employee attrition. By enhancing minority-class representation, SMOTE significantly improves recall and overall prediction performance, making the model more effective in identifying employees at risk of leaving [15].

3.3 Model Building

The model building phase involves constructing an efficient and robust machine learning pipeline for predicting employee attrition. The proposed system integrates multiple stages, including data preprocessing, class imbalance handling, feature optimization, ensemble classification, and evaluation. Initially, the preprocessed dataset is divided into training and testing sets using a stratified approach to maintain the distribution of attrition classes. To address class imbalance, the Synthetic Minority Oversampling Technique (SMOTE) is applied to generate synthetic samples for the minority class, thereby improving the model's ability to detect attrition cases effectively [15].

Following preprocessing, feature optimization is performed using a Genetic Algorithm (GA), which selects the most relevant subset of features by iteratively applying selection, crossover, and mutation operations. This evolutionary approach helps in reducing dimensionality, eliminating redundant features, and improving model efficiency and performance [18].

The optimized feature set is then used to train an ensemble classification model combining LightGBM and CatBoost classifiers. LightGBM efficiently handles large-scale structured data and captures complex feature interactions, while CatBoost effectively processes categorical features and reduces prediction bias [8], [6]. A soft-voting mechanism is employed to aggregate predictions from both models, enhancing predictive robustness and generalization capability.




Figure 3.1: Proposed Architecture

3.3.1 Genetic Algorithm-Based Feature Selection

The proposed system utilizes a Genetic Algorithm (GA) to perform feature selection and optimize the input feature set. GA is an evolutionary optimization technique inspired by the principles of natural selection, where each candidate solution represents a subset of features encoded as a binary chromosome [18]. The algorithm begins with an initial population of randomly generated feature subsets and evaluates their performance using a fitness function based on model performance metrics such as accuracy or F1-score.

Through iterative processes such as selection, crossover, and mutation, the GA evolves toward an optimal feature subset. Tournament selection is employed to choose high-performing individuals, while crossover combines feature subsets from parent solutions, and mutation introduces randomness to maintain diversity in the population. These operations help prevent premature convergence and improve the search for a global optimum [14].

The evolutionary process continues for multiple generations until convergence is achieved or a stopping criterion is met. The final output is an optimized subset of features that significantly reduces dimensionality, eliminates redundant or irrelevant attributes, and enhances model efficiency. This approach not only improves predictive performance but also reduces computational complexity, making it suitable for large-scale HR analytics datasets [13].

3.3.2 Ensemble Classification Model

The classification stage employs an ensemble learning approach by combining LightGBM and CatBoost classifiers. LightGBM is a gradient boosting algorithm known for its efficiency and ability to handle large-scale structured data [8], while CatBoost is specifically designed to handle categorical features effectively and reduce prediction bias [6]. By combining these two models, the system leverages their complementary strengths to improve overall performance.

A soft-voting mechanism is used to aggregate the predicted probabilities from both classifiers, where each model contributes to the final decision based on its confidence. This ensemble approach improves predictive stability, reduces variance, and enhances the model's ability to capture complex patterns in employee data. Additionally, ensemble methods are known to outperform single-model approaches by combining diverse learning behaviors, resulting in more robust and reliable predictions [12].

Furthermore, the integration of LightGBM and CatBoost in a hybrid ensemble framework enhances the model's generalization capability by combining their individual strengths in handling structured and categorical data. The soft-voting strategy ensures that the final prediction is more balanced and less sensitive to individual model errors. This results in improved classification performance, especially in complex HR datasets, making the model more effective and reliable for real-world employee attrition prediction tasks.

3.3.3 Prediction Layer

After training, the ensemble model generates predictions for employee attrition based on the optimized feature set. The prediction layer outputs whether an employee is likely to stay or leave the organization, along with associated probability scores. These probability values indicate the confidence of the model in its predictions and help in ranking employees based on their risk of attrition.

The use of ensemble-based probability estimation improves prediction reliability by combining outputs from multiple models, leading to more stable and accurate results compared to single-model approaches [12]. Additionally, probability-based predictions support threshold-based decision-making, allowing organizations to define risk levels and take proactive measures for employee retention.

These predictions play a crucial role in supporting data-driven HR decision-making by enabling early identification of high-risk employees. As a result, organizations can implement targeted interventions such as employee engagement strategies,

performance improvements, or policy adjustments to reduce attrition rates and improve workforce stability.

3.3.4 Performance Evaluation

The performance of the proposed model is evaluated using standard classification metrics such as accuracy, precision, recall, F1-score, and ROC-AUC. These metrics provide a comprehensive understanding of the model's effectiveness, particularly in handling imbalanced datasets. Accuracy measures the overall correctness of predictions, while precision and recall evaluate the model's ability to correctly identify attrition cases without generating excessive false positives or false negatives. The F1-score provides a balanced measure by combining precision and recall, making it especially useful for imbalanced classification problems [13].

Additionally, ROC-AUC (Receiver Operating Characteristic – Area Under Curve) is used to assess the model's ability to distinguish between employees who stay and those who leave across different threshold values [13]. A higher ROC-AUC score indicates better classification performance and robustness. These evaluation metrics collectively ensure that the model not only achieves high predictive accuracy but also maintains reliability and consistency in real-world HR scenarios.

3.3.5 Explainability using SHAP

To enhance interpretability, SHAP (SHapley Additive exPlanations) is used to analyze feature contributions. SHAP values quantify the impact of each feature on the prediction outcome by assigning importance scores based on their contribution to the model's output [9]. This allows the identification of key factors influencing employee attrition at both global and individual levels.

By providing clear and interpretable insights, SHAP enables HR professionals to understand the reasoning behind model predictions and make informed decisions. This transparency improves trust in the system and supports effective, data-driven strategies for employee retention, making the model more suitable for real-world HR analytics applications.

3.3.6 Loss Function: Binary Cross-Entropy

The proposed system uses Binary Cross-Entropy (BCE) as the loss function, which is well-suited for binary classification problems such as employee attrition prediction. In this task, the model predicts whether an employee will leave the organization (Attrition = 1) or stay (Attrition = 0). The model outputs a probability score representing the likelihood of attrition, which is compared with the actual class label.

Binary Cross-Entropy measures the difference between the predicted probability and the true label by penalizing incorrect predictions. It assigns higher loss values when the predicted probability deviates significantly from the actual outcome, thereby guiding the model toward better predictions. Minimizing this loss function encourages the model to assign higher probabilities to the correct class and lower probabilities to the incorrect class [12].

As training progresses, a reduction in loss indicates improved model performance and better alignment between predicted and actual values. This makes Binary Cross-Entropy an effective choice for optimizing classification accuracy and ensuring reliable prediction outcomes in employee attrition analysis.

3.4 Model Validation

Model validation is a crucial step to ensure that the proposed hybrid machine learning model performs effectively on unseen employee data. In this study, validation is carried out to assess the generalization capability, robustness, and reliability of the attrition prediction framework [12].

The dataset is divided into training and testing subsets using a stratified approach, typically in an 80:20 ratio. The training data is used to build and optimize the model, while the test data is reserved for evaluating performance on unseen instances. Stratification ensures that the class distribution of attrition and non-attrition cases remains consistent across both subsets, which is essential for imbalanced classification problems [15].

During training, validation techniques are applied to monitor model performance and prevent overfitting. The incorporation of SMOTE and Genetic Algorithm-based feature selection further enhances model stability by improving class balance and selecting relevant features [15][18]. Additionally, ensemble learning through LightGBM and CatBoost contributes to improved generalization by combining the strengths of multiple classifiers [8][6].

By validating the model on unseen data and maintaining controlled experimental conditions, the proposed system achieves consistent and reliable performance. This ensures that the model can effectively predict employee attrition in real-world HR analytics scenarios.

3.5 Evaluation Metrics

To evaluate the performance of the proposed hybrid model, several standard classification metrics are used, including Accuracy, Precision, Recall, F1-score, and Area Under the ROC Curve . These metrics provide a comprehensive understanding of the model's predictive capability, especially in imbalanced datasets [13].

Accuracy measures the overall correctness of the model by calculating the proportion of correctly predicted instances out of all predictions. However, in attrition prediction, accuracy alone may not be sufficient due to class imbalance, as it may favor the majority class.

$$Accuracy = \frac{TP+TN}{FP+FN+TP+TN}$$

Precision indicates the proportion of correctly predicted attrition cases among all predicted attrition instances. It reflects the model's ability to avoid false positives and ensures that predicted attrition cases are reliable.

$$Precision = \frac{TP}{TP+FP}$$

Recall measures the proportion of actual attrition cases that are correctly identified by the model. This metric is particularly important in HR analytics, as missing true attrition cases can lead to significant organizational impact.

$$Recall = \frac{TP}{TP+FN}$$

F1-score is the harmonic mean of precision and recall, providing a balanced evaluation of the model's performance. It is especially useful when dealing with imbalanced datasets, as it considers both false positives and false negatives [13].

$$F1 - Score = 2 * \frac{Precision * Recall}{Precision + Recall}$$

AUC (Area Under the ROC Curve) evaluates the model's ability to distinguish between attrition and retention classes across different threshold values. A higher AUC indicates better discriminative performance and overall model robustness [13].

$$PR-AUC = \int Precision d(Recall)$$

These evaluation metrics collectively ensure a comprehensive assessment of the proposed model, highlighting its effectiveness in accurately identifying employee attrition while maintaining balanced, reliable, and robust predictions across different evaluation scenarios [13].

3.6 Feasibility Study

3.6.1 Economic Feasibility

The proposed employee attrition prediction system is economically feasible as it is developed using open-source technologies such as Python, scikit-learn, LightGBM, CatBoost, and SHAP, eliminating any licensing costs. The model training and evaluation are performed using platforms like Google Colab, which provide free computational resources, reducing infrastructure expenses. Additionally, the system does not require specialized hardware and can be executed on standard computing environments. Since the dataset used is publicly available, there are no data acquisition costs, making the overall implementation cost-effective and suitable for organizations of varying scales [8][6].

Furthermore, the use of efficient machine learning models reduces computational time and resource consumption. The absence of expensive deployment requirements makes the system affordable for both small and large organizations. This ensures long-term sustainability and low maintenance cost.

3.6.2 Operational Feasibility

The proposed system is operationally feasible and can be easily integrated into existing HR analytics platforms or decision-support systems. It provides automated predictions of employee attrition based on structured workforce data, reducing manual analysis efforts. The inclusion of SHAP-based explainability enhances usability by allowing HR professionals to understand the factors influencing predictions, thereby improving trust and decision-making [9].

The system is user-friendly, requires minimal human intervention, and can be regularly updated with new employee data, ensuring continuous improvement and adaptability in dynamic organizational environments. This makes it practical for real-world deployment in HR analytics and workforce management systems.

3.6.3 Technical Feasibility

The technical feasibility of the proposed system is high, as it leverages well-established machine learning techniques and widely used libraries. The implementation is carried out using Python along with efficient frameworks such as LightGBM and CatBoost for model building, and SHAP for interpretability [8][6][9]. These tools are well-supported, scalable, and suitable for handling structured HR datasets.

Data preprocessing techniques, including normalization, categorical encoding, and SMOTE-based imbalance handling, are computationally efficient and commonly used in machine learning workflows [12][15]. The application of Genetic Algorithm (GA) for feature selection further enhances model performance by selecting relevant features while maintaining manageable computational complexity [18].

Moreover, the system is capable of being deployed in cloud-based environments, ensuring scalability and accessibility for large-scale organizational use. The modular design of the framework allows easy integration with existing HR systems and supports future enhancements without significant technical modifications. This ensures long-term usability and adaptability of the proposed solution.

Chapter 4

System Requirements

4. SYSTEM REQUIREMENTS

4.1 Functional Requirements

The functional requirements define the core operations that the proposed employee attrition prediction system must perform. The system is designed to analyze structured HR data and predict employee attrition using a hybrid machine learning framework.

The system should be capable of loading and preprocessing HR datasets, including handling missing values, encoding categorical variables, and normalizing numerical features to ensure consistent and high-quality input data [12]. It must support class imbalance handling using the Synthetic Minority Oversampling Technique (SMOTE) to improve minority-class representation and enhance prediction performance [15]. Additionally, the system should implement Genetic Algorithm (GA)-based feature selection to identify the most relevant attributes and reduce redundancy in the dataset [18].

The system should build and train a hybrid ensemble model combining LightGBM and CatBoost classifiers using a soft-voting mechanism, enabling improved predictive robustness and accuracy [8][6]. It must generate predictions indicating whether an employee is likely to leave the organization, along with probability scores.

Furthermore, the system should evaluate model performance using metrics such as accuracy, precision, recall, F1-score, and AUC to ensure reliable and balanced predictions [13]. The system should also provide interpretability through SHAP, enabling users to understand feature contributions influencing attrition predictions and improving trust in the model [9].

Finally, the system should generate outputs such as prediction results, performance reports, and visualizations (e.g., ROC curves, confusion matrix, feature importance) to support effective decision-making in HR analytics.

4.2 Non-Functional Requirements

Non-functional requirements describe the quality attributes of the system, focusing on performance, usability, scalability, and reliability.

The system should deliver accurate and efficient predictions while handling moderately large HR datasets. It must be scalable to accommodate increasing data volumes and adaptable to different organizational environments. The use of efficient machine learning frameworks such as LightGBM and CatBoost ensures faster processing and improved performance during training and prediction stages [8][6].

Usability is an important aspect, and the system should be easy to use and understandable for HR professionals, even with limited technical knowledge. The inclusion of SHAP-based explainability enhances transparency and trust by clearly presenting feature contributions influencing predictions [9].

The system should be reliable and produce consistent results across different runs. It must also be maintainable, with a modular design that allows easy updates, integration of new data, and future enhancements. Additionally, the system should ensure data integrity and support secure handling of employee data, making it suitable for real-world HR analytics applications.

4.2.1 Software Requirements

The proposed system requires a software environment that supports machine learning, data processing, and model interpretability.

Operating System: Windows 10 / 11 (64-bit)

Programming Language: Python 3.8 or higher

Development Environment: Google Colab / Jupyter Notebook

Frameworks & Libraries:

Machine Learning: LightGBM, CatBoost, scikit-learn[8][6]

Data Processing: NumPy, Pandas

Visualization: Matplotlib, Seaborn

Imbalance Handling: imbalanced-learn (SMOTE)[15]

Explainability: SHAP[9]

These tools provide efficient implementation of preprocessing, feature selection, model training, and evaluation processes, enabling the development of a robust and scalable employee attrition prediction system..

4.2.2 Hardware Requirements

The hardware requirements for the proposed system are minimal and can be satisfied using standard computing resources.

Memory (RAM): Minimum 4 GB (Recommended: 8 GB or above)

Processor: Intel Core i5 or equivalent

Processor Speed: 2.0 GHz or higher

Graphics: Optional GPU (NVIDIA CUDA support preferred for training)

Monitor: 15” display with minimum 720p resolution

Storage: Minimum 10–20 GB of free space for datasets and model checkpoints

A GPU is not mandatory but can be beneficial for faster computation when working with large datasets. Cloud platforms such as Google Colab can also be used to access GPU resources without dedicated hardware. A stable internet connection is recommended for cloud-based execution and dataset access.

This configuration ensures smooth data processing, efficient model training, and reliable system performance in real-world environments.

Chapter 5

Design

5.DESIGN

5.1 System design




Figure 5.1: Block Diagram of Proposed Hybrid Employee Attrition Prediction

The proposed system is designed as a multi-layered hybrid machine learning architecture that integrates data preprocessing, imbalance handling, feature optimization, ensemble learning, and explainability to predict employee attrition effectively. This layered structure ensures a systematic flow of data from raw input to final prediction and interpretation, improving both performance and usability.

The system is organized into the following layers: Each layer performs a specific function, ensuring modularity and efficient data processing. This structured design improves scalability and allows easy integration of additional techniques or future enhancements. It also ensures better accuracy, reliability, and interpretability of the overall system.

Furthermore, the layered architecture enables clear separation of responsibilities across different modules, simplifying system development and maintenance. It also

facilitates easier debugging and performance optimization by isolating individual components. This design enhances the overall flexibility and robustness of the proposed system.

1. Data Source Layer

The process begins with the HR dataset collected from structured sources such as employee databases or publicly available datasets. This dataset includes attributes such as age, job role, salary, job satisfaction, work-life balance, and performance indicators. These features capture both personal and professional aspects of employees, which are essential for identifying patterns related to attrition.

This layer acts as the foundation of the system, as the quality and diversity of input data directly influence model performance and prediction accuracy.

2. Data Processing Layer

In this layer, raw data undergoes preprocessing to ensure quality, consistency, and usability. This includes handling missing values, removing duplicate records, encoding categorical variables, normalization, feature scaling, and elimination of irrelevant features [12].

These preprocessing steps reduce noise, standardize feature values, and prepare the dataset for efficient model training. Proper data processing improves learning capability and ensures that the model can generalize well to unseen data.

3. SMOTE Oversampling Layer

To address class imbalance, the Synthetic Minority Oversampling Technique (SMOTE) is applied. In HR datasets, attrition cases are usually fewer than retention cases, which can bias the model. SMOTE generates synthetic samples for the minority class, improving class distribution and enabling the model to learn balanced decision boundaries [15].

This layer significantly enhances the model's ability to detect employees who are likely to leave, thereby improving recall and overall prediction reliability.

4. Data Split Layer

The processed dataset is divided into training and testing sets using a stratified train–test split approach. This ensures that both subsets maintain the same proportion of attrition and non-attrition cases [15].

The training set is used to build the model, while the testing set evaluates performance on unseen data. This separation prevents overfitting and ensures realistic performance estimation.

5. Feature Optimization Layer (Genetic Algorithm)

A Genetic Algorithm (GA) is used to select the most relevant features from the dataset. Each feature subset is represented as a chromosome, and the algorithm iteratively applies selection, crossover, and mutation operations to find the optimal feature combination [18].

This process reduces dimensionality, eliminates redundant features, and improves computational efficiency. It also enhances model accuracy by focusing only on the most significant attributes influencing employee attrition.

6. Ensemble Classifier Model Layer

The optimized features are fed into two advanced machine learning models: LightGBM and CatBoost. LightGBM efficiently handles large-scale structured data, while CatBoost effectively processes categorical features and reduces prediction bias [8][6].

These models are combined using a soft voting mechanism, where predicted probabilities from both models are aggregated to produce the final output. This ensemble approach improves prediction accuracy, stability, and robustness by leveraging the strengths of multiple algorithms.

7. Prediction Layer

The ensemble model generates predictions indicating whether an employee is likely to leave or stay. Along with classification output, the model provides probability scores that represent the likelihood of attrition.

These probability-based predictions help HR professionals identify high-risk employees and prioritize retention strategies accordingly. This layer plays a key role in enabling proactive decision-making.

8. Results & Interpretation Layer

a) Model Interpretability (SHAP)

SHAP is used to interpret model predictions by analyzing feature contributions. It assigns importance scores to each feature, showing how they influence the prediction outcome [9].

This enhances transparency and allows HR professionals to understand the reasoning behind predictions, increasing trust in the system .

b) Evaluation & Output

The final stage evaluates model performance using metrics such as accuracy, precision, recall, F1-score, and ROC-AUC [13]. These metrics provide a comprehensive understanding of model effectiveness.

The system generates outputs including:

- Prediction results (Stay/Leave)
- Probability scores
- Performance reports
- Visualizations (ROC curve, confusion matrix, feature importance)

These outputs support data-driven HR decision-making and help organizations implement effective employee retention strategies.

The ensemble model generates predictions indicating whether an employee is likely to leave or stay. Along with classification output, the model provides probability scores that represent the likelihood of attrition.

These probability-based predictions help HR professionals identify high-risk employees and prioritize retention strategies accordingly. This layer plays a key role in enabling proactive decision-making.

5.2 UML Diagrams

5.2.1 Class Diagram




Figure 5.2.1: UML Class Diagram of Employee Attrition Prediction System

The class diagram represents the structural architecture of the proposed employee attrition prediction system. It illustrates the various system components in terms of classes, along with their attributes, methods, and relationships. The User class is responsible for interacting with the system by uploading datasets and viewing prediction results. The Dataset class manages data loading and validation processes. The Preprocessor class performs data cleaning, encoding of categorical variables, and feature scaling to prepare the dataset for model training [12].

The ModelTrainer class is used to train machine learning models such as LightGBM and CatBoost [8][6]. The Predictor class utilizes the trained model to generate attrition predictions. Finally, the Evaluator class assesses model performance using evaluation metrics [13]. The relationships between these classes ensure a modular, scalable, and well-structured system design, facilitating easy maintenance and extension of the system [14].

5.2.2 Sequence Diagram




Figure 5.2.2: UML Sequence Diagram for System Interaction Flow

The sequence diagram illustrates the dynamic interaction between system components over time. It shows the step-by-step flow of operations from user input to final prediction output. Initially, the User uploads the dataset, which is received by the Dataset module. The dataset is then loaded and passed to the Preprocessor, where data cleaning and encoding operations are performed [12]. The processed data is forwarded to the Model, where training is carried out using machine learning algorithms such as LightGBM and CatBoost [8][6].

Once the model is trained, the Predictor generates predictions based on the processed data. Finally, the prediction results are returned to the User. This diagram clearly demonstrates the sequence of operations and the communication between different modules in the system, ensuring efficient data processing and prediction workflow [14].

5.2.3 Use Case Diagram




Figure 5.2.3: UML Use Case Diagram of Employee Attrition Prediction System

The use case diagram represents the interaction between the user and the system, highlighting the functionalities provided. The primary actor in the system is the User (HR/Admin), who interacts with the employee attrition prediction system. The user can perform several operations such as:

- Uploading the dataset
- Preprocessing data
- Training the model
- Predicting employee attrition
- Viewing prediction results

This diagram provides a high-level view of system functionalities and helps in understanding user requirements and system behavior.

It also simplifies system analysis by clearly defining user interactions and system boundaries. Additionally, it assists in identifying key functionalities required for effective implementation and user experience [14][11].

5.2.4 Workflow Diagram




Figure 5.2.4: Workflow Diagram of Proposed Hybrid Machine Learning Model

The workflow diagram illustrates the complete operational pipeline of the proposed hybrid machine learning system for employee attrition prediction. The process begins with dataset collection, followed by data preprocessing, which includes handling missing values, encoding categorical variables, and normalization [12]. Next, SMOTE is applied to address class imbalance [15].

The dataset is then divided using a train-test split approach. Genetic Algorithm (GA) is applied for feature selection to identify the most relevant features [18]. The selected features are used to train LightGBM and CatBoost models, which are combined using a soft voting mechanism [8][6]. The system then performs prediction, followed by evaluation using performance metrics [13]. Finally, SHAP explainability is applied to interpret model predictions [9].

Chapter 6

Implementation

6. IMPLEMENTATION

6.1 Implementation

STEP 0: INSTALL

```
!pip install lightgbm catboost imbalanced-learn shap scikit-learn pandas numpy  
matplotlib seaborn -q
```

STEP 1 — IMPORTS

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import shap
```

```
import warnings, random, time
```

```
warnings.filterwarnings('ignore')
```

```
from google.colab import files
```

```
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
```

```
from sklearn.model_selection import StratifiedKFold, cross_val_score,  
train_test_split
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,  
roc_auc_score, confusion_matrix, ConfusionMatrixDisplay
```

```
from imblearn.over_sampling import SMOTE
```

```
import lightgbm as lgb
```

```

from catboost import CatBoostClassifier

from sklearn.ensemble import VotingClassifier

SEED = 42

np.random.seed(SEED)

random.seed(SEED)

# STEP 2: LOAD DATA

df = pd.read_csv('/content/HR_Analytics.csv', encoding='utf-8-sig')

df = df.drop_duplicates().reset_index(drop=True)

# Encode target

df['Attrition'] = (df['Attrition'] == 'Yes').astype(int)

# Drop unnecessary columns

drop_cols = ['EmpID', 'EmployeeNumber', 'EmployeeCount', 'StandardHours',
'Over18', 'AgeGroup', 'SalarySlab']

df.drop(columns=[c for c in drop_cols if c in df.columns], inplace=True)

# Handle missing values

for col in df.select_dtypes(include=np.number).columns:

    df[col].fillna(df[col].mean(), inplace=True)

le = LabelEncoder()    # Encode categorical features

for col in df.select_dtypes(include='object').columns:

    df[col] = le.fit_transform(df[col].astype(str))

```



```

# Split features and target

X = df.drop('Attrition', axis=1)

y = df['Attrition'].values

feature_names = list(X.columns)

scaler = MinMaxScaler() # Scaling

X[X.columns] = scaler.fit_transform(X)

X = X.values

# STEP 3: SMOTE

smote = SMOTE(random_state=SEED)

X_res, y_res = smote.fit_resample(X, y)

# Train/Test Split

X_train, X_test, y_train, y_test = train_test_split(

    X_res, y_res, test_size=0.2, stratify=y_res, random_state=SEED)

# STEP 4: GENETIC ALGORITHM (Feature Selection)

POP_SIZE, N_GEN, MUT_RATE, ELITE = 30, 20, 0.1, 2

n_features = X_train.shape[1]

def fitness_fn(chromosome):

    selected = np.where(chromosome == 1)[0]

    if len(selected) < 5:

        return 0.0

```

```

model = lgb.LGBMClassifier(

    n_estimators=150, learning_rate=0.05,

    random_state=SEED, class_weight='balanced', verbosity=-1)

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=SEED)

return cross_val_score(model, X_train[:, selected], y_train, cv=cv,
scoring='f1').mean()

def mutate(ch):

    for i in range(len(ch)):

        if random.random() < MUT_RATE:

            ch[i] = 1 - ch[i]

    return ch

population = np.random.randint(0, 2, size=(POP_SIZE, n_features))

best = population[0]

for gen in range(N_GEN):

    scores = np.array([fitness_fn(ind) for ind in population])

    best = population[np.argmax(scores)]

    new_pop = population[np.argsort(scores)[-ELITE:]].tolist()

    while len(new_pop) < POP_SIZE:

        p = population[np.random.randint(POP_SIZE)]

        new_pop.append(mutate(p.copy()))

```

```

population = np.array(new_pop)

selected_idx = np.where(best == 1)[0]

selected_names = [feature_names[i] for i in selected_idx]

X_train_ga = X_train[:, selected_idx]

X_test_ga = X_test[:, selected_idx]

print(f'Selected {len(selected_names)} features")

# STEP 5 — PROPOSED MODEL (Ensemble)

lgbm = lgb.LGBMClassifier(

    n_estimators=700, learning_rate=0.02,

    num_leaves=63, max_depth=7,

    subsample=0.9, colsample_bytree=0.9,

    class_weight='balanced', random_state=SEED, verbosity=-1)

cat = CatBoostClassifier(iterations=700, learning_rate=0.02, depth=7, verbose=0,

    random_state=SEED, auto_class_weights='Balanced')

model = VotingClassifier(estimators=[('lgbm', lgbm), ('cat', cat)],

    voting='soft', weights=[2,1])

print("Training Proposed Model")

model.fit(X_train_ga, y_train)

# STEP 6: EVALUATION

y_pred = model.predict(X_test_ga)

```

```

y_prob = model.predict_proba(X_test_ga)[:,-1]

print("\n PROPOSED MODEL RESULTS ")

print("Accuracy :", round(accuracy_score(y_test, y_pred),2))

print("Precision:", round(precision_score(y_test, y_pred),2))

print("Recall  :", round(recall_score(y_test, y_pred),2))

print("F1 Score :", round(f1_score(y_test, y_pred),2))

print("AUC    :", round(roc_auc_score(y_test, y_prob),2))

# STEP 7: CONFUSION MATRIX

ConfusionMatrixDisplay(confusion_matrix(y_test, y_pred),

                        display_labels=['Stay','Leave']).plot()

plt.title("Proposed Model Confusion Matrix")

plt.show()

# STEP 8: SHAP (Explainability)

lgbm_model = model.named_estimators_['lgbm']

explainer = shap.TreeExplainer(lgbm_model)

X_test_df = pd.DataFrame(X_test_ga, columns=selected_names)

shap_values = explainer.shap_values(X_test_df)

sv = shap_values[1] if isinstance(shap_values, list) else shap_values

shap.summary_plot(sv, X_test_df, plot_type="bar")

```

Chapter 7

Results

7. RESULTS

The results section evaluates the performance of the proposed hybrid machine learning framework for employee attrition prediction. The system was implemented using the HR Analytics dataset and includes preprocessing, SMOTE-based imbalance handling, Genetic Algorithm (GA) feature selection, and an ensemble classification approach.

The dataset initially contained 1473 records with 30 features, and significant class imbalance was observed between retention and attrition instances. After applying SMOTE, the dataset was balanced, improving the model's ability to learn minority-class patterns effectively [15]. The dataset was then split into training and testing sets using an 80:20 stratified strategy to ensure unbiased and reliable evaluation [15].

The proposed model integrates LightGBM and CatBoost using a soft voting ensemble, enabling improved predictive stability and robustness compared to single-model approaches. By combining the strengths of both models, the system effectively captures complex feature interactions and enhances overall prediction performance [8][6].

The evaluation results demonstrate that the hybrid model achieves high accuracy, precision, recall, and F1-score, indicating strong performance in both majority and minority class predictions. The ROC-AUC score further confirms the model's strong discriminative ability across different classification thresholds [13].

Additionally, SHAP-based explainability provides insights into important features influencing employee attrition, enabling HR professionals to make informed decisions [9]. Overall, the proposed framework outperforms traditional approaches and offers a reliable solution for real-world HR analytics applications. The results highlight the effectiveness of integrating multiple techniques into a unified framework. Furthermore, the model demonstrates strong generalization capability across unseen data.

7.1 Model Results




Figure 7.1: ROC Curve Comparison

The proposed hybrid model achieved strong predictive performance on the test dataset. The integration of GA-based feature selection and ensemble learning significantly improved the model's ability to detect attrition cases effectively by selecting the most relevant features and reducing redundancy [18]. Compared to the existing GA + LightGBM model, the proposed system shows improved predictive capability, particularly in handling class imbalance and capturing complex relationships in HR data. The use of SMOTE further enhances minority-class detection, leading to better recall and balanced performance [15].

Moreover, the ensemble approach combining LightGBM and CatBoost ensures better generalization and stability by leveraging the strengths of multiple models [8][6]. This results in reduced variance and more consistent predictions across different data samples. The incorporation of SHAP-based explainability also adds value by providing clear insights into feature contributions, enabling better understanding and trust in the model's predictions [9]. Overall, the proposed system demonstrates superior performance, robustness, and practical applicability for real-world employee attrition prediction.

7.2 Results Comparison

Table 7.1: Comparison of Classification Metrics for Existing and Proposed Models

Metric	Existing GA+LightGBM Model	Proposed GA+LightGBM+CatBoost Model
Accuracy	92%	93%
Precision	89%	96%
Recall	85%	90%
F1-Score	87%	93%

A comparative analysis was conducted between:

- Existing Model: GA + LightGBM
- Proposed Model: GA + LightGBM + CatBoost
- Baseline Models (Logistic Regression, Random Forest, SVM, etc.).

The proposed model consistently outperforms other models across multiple evaluation metrics such as accuracy, precision, recall, F1-score, and ROC-AUC, demonstrating its superior predictive capability [13]. This improvement is primarily due to:

- Effective feature selection using Genetic Algorithm
- Balanced learning using SMOTE
- Complementary learning of LightGBM and CatBoost

In comparison to baseline models, traditional algorithms such as Logistic Regression and SVM show limited performance due to their inability to effectively model nonlinear relationships in high-dimensional HR data. Although Random Forest provides better results than basic models, it still lacks the optimization and ensemble diversity achieved in the proposed approach.

The results confirm that hybrid ensemble learning provides better predictive reliability, robustness, and generalization compared to single-model approaches. This

makes the proposed system more suitable for real-world employee attrition prediction and HR decision-making scenarios.




Figure 7.2: Comparison of Proposed and Existing Models




Figure 7.3: Comparative Performance of All Machine Learning Models

7.3 ROC Curve Analysis

The ROC curve analysis illustrates the classification capability of all models by comparing the trade-off between true positive rate (TPR) and false positive rate (FPR) across different threshold values. The proposed model achieves the highest Area Under Curve (AUC), indicating superior class separation between attrition and retention cases [13].

The curve shows that the proposed model consistently maintains higher true positive rates while minimizing false positives across various thresholds, demonstrating its robustness and effectiveness in classification. This is particularly important in employee attrition prediction, where accurately identifying employees at risk is critical.

Compared to baseline and existing models, the hybrid ensemble approach provides a smoother and more dominant ROC curve, reflecting better discriminative power. The integration of SMOTE and ensemble learning contributes to improved sensitivity toward minority-class instances, enhancing overall model performance [15][8].

Overall, the ROC curve analysis confirms that the proposed model offers strong predictive capability, better generalization, and reliable performance in real-world HR analytics scenarios.

7.4 Confusion Matrix Analysis

The confusion matrix provides a detailed view of prediction outcomes by comparing actual and predicted values. It includes:

- True Positives (TP): correctly predicted attrition cases
- True Negatives (TN): correctly predicted retention cases
- False Positives (FP): incorrect attrition predictions
- False Negatives (FN): missed attrition cases

The proposed model shows fewer misclassifications compared to the existing model, particularly reducing false negatives, which is critical in attrition prediction since

undetected attrition cases can lead to organizational loss. This improvement highlights the model's strong recall performance and its ability to accurately identify employees at risk [13].

Additionally, the balanced distribution of predictions indicates that the model effectively handles class imbalance and avoids bias toward the majority class. The integration of SMOTE and ensemble learning contributes to improved classification accuracy and reduced error rates [15][8]. Overall, the confusion matrix analysis demonstrates that the proposed hybrid model provides reliable and consistent prediction outcomes, making it suitable for real-world HR analytics applications..




Figure 7.4 : Confusion Matrix

7.5 Feature Importance Analysis

Feature importance analysis using LightGBM highlights the most influential factors contributing to employee attrition. The model identifies key attributes such as job satisfaction, overtime, monthly income, and work environment as major contributors to prediction outcomes [8].

This analysis helps in understanding which features significantly impact employee decisions to leave the organization. By ranking features based on their importance scores, the model provides insights into critical factors affecting attrition, enabling HR

professionals to focus on key areas such as employee satisfaction, workload management, and compensation.

Additionally, feature importance analysis improves model interpretability by revealing how different variables influence predictions. When combined with SHAP-based explanations, it provides both global and local insights into feature contributions [9]. Overall, this analysis supports data-driven decision-making and helps organizations implement targeted strategies to reduce employee attrition.




Figure 7.5 : Feature Importance

7.6 Model Explainability using SHAP

SHAP (SHapley Additive Explanations) is used to interpret model predictions by providing detailed insights into feature contributions. It offers both global and local explanations, enhancing the transparency of the proposed system. Global explanation highlights the overall importance of features across the dataset, while local explanation focuses on individual predictions, showing how each feature influences a specific outcome [9].

SHAP visualizations clearly demonstrate how features contribute positively or negatively to employee attrition prediction. For instance, factors such as overtime or low job satisfaction may increase attrition probability, while higher income or better work-life balance may reduce it. These insights help in understanding the decision-making process of the model.

By improving interpretability, SHAP enables HR professionals to trust the model's predictions and make informed decisions. This transparency makes the system more reliable and suitable for real-world HR analytics applications, where explainability is a critical requirement.




Figure 7.6 : SHAP Feature Importance

Chapter 8

Social Impact

8. SOCIAL IMPACT

The development of an intelligent employee attrition prediction system using machine learning has significant social and organizational impact. In modern workplaces, employee turnover is a critical challenge that affects productivity, organizational stability, and employee morale. Traditional methods of identifying at-risk employees rely heavily on manual observation and basic statistical analysis, which are often inefficient and unable to capture complex patterns in workforce data. This leads to delayed decision-making and ineffective retention strategies.

The proposed system addresses these challenges by providing a data-driven and predictive approach to employee attrition. By integrating advanced techniques such as SMOTE for handling class imbalance, Genetic Algorithm for feature selection, and ensemble learning using LightGBM and CatBoost, the system enables organizations to accurately identify employees who are likely to leave. This allows HR professionals to take proactive measures such as improving work conditions, offering incentives, or providing career development opportunities, thereby reducing attrition rates and enhancing employee satisfaction.

From a societal perspective, the system contributes to improved job stability and employee well-being. Early identification of attrition risk helps organizations create a supportive work environment, reducing stress and uncertainty among employees. By promoting fair and informed decision-making, the system also helps minimize bias in HR processes. The inclusion of SHAP-based explainability further enhances transparency by allowing HR managers to understand the key factors influencing predictions, fostering trust in AI-driven decisions.

The proposed system also supports organizational efficiency and economic growth. Reducing employee turnover minimizes recruitment and training costs while maintaining workforce continuity. This is particularly beneficial for large organizations and industries where high attrition rates can significantly impact operational performance. Additionally, the system can be integrated into HR analytics

platforms, enabling scalable and real-time decision support across different organizational levels.

Furthermore, the project promotes responsible and ethical use of artificial intelligence in human resource management. By focusing on interpretability and transparency, the system ensures that predictions are explainable and can be validated by decision-makers. This aligns with the growing need for trustworthy AI systems in real-world applications.

Additionally, the system can be extended to support personalized employee engagement strategies by analyzing individual-level insights and behavioral patterns. Organizations can use these insights to design targeted retention policies, improve workplace culture, and enhance employee experience. Such data-driven personalization not only strengthens employee–organization relationships but also contributes to long-term workforce sustainability and organizational success.

Moreover, the system has the potential to be integrated with real-time HR management systems, enabling continuous monitoring of employee behavior and timely intervention strategies. This real-time capability allows organizations to respond quickly to emerging attrition risks, making workforce management more proactive rather than reactive. Such advancements further strengthen the role of intelligent systems in building adaptive, resilient, and future-ready organizations.

Overall, the proposed employee attrition prediction system contributes to creating a more stable, efficient, and employee-centric workplace. By combining advanced machine learning techniques with explainable AI, the system enhances decision-making, improves employee retention, and supports the development of smarter and more sustainable organizational practices.

Chapter 9

Conclusion & Future Work

9. CONCLUSION & FUTURE WORK

In conclusion, an efficient and interpretable employee attrition prediction system was successfully developed using a hybrid machine learning approach. The primary objective of this work was to accurately predict employee attrition by addressing key challenges such as class imbalance, feature redundancy, and lack of interpretability in traditional models. This objective was achieved by designing a structured pipeline that includes data preprocessing, SMOTE-based imbalance handling, Genetic Algorithm-based feature selection, and an ensemble model combining LightGBM and CatBoost.

The system was implemented using Python in a cloud-based environment, ensuring computational efficiency and scalability. During preprocessing, the dataset was cleaned, encoded, and normalized to improve data quality and model performance. Feature optimization using Genetic Algorithm significantly reduced redundancy and improved model robustness. The integration of ensemble learning enabled the system to capture complex patterns in HR data, resulting in reliable and consistent predictions.

The proposed model demonstrated strong performance and improved capability in identifying employees at risk of leaving the organization. The inclusion of SHAP-based explainability provided clear insights into feature contributions, enhancing transparency and trust in the model. This makes the system suitable for real-world HR analytics applications where both accuracy and interpretability are essential.

The developed system meets both technical and practical objectives and has strong potential for deployment in organizational environments. It enables proactive decision-making by helping HR professionals identify attrition risks early and take preventive actions, ultimately improving employee retention and organizational stability.

Despite its effectiveness, there is scope for further improvement. Future work can focus on extending the model to handle real-time streaming data and dynamic

workforce changes. Incorporating deep learning techniques or hybrid architectures may further enhance predictive performance. Additionally, integrating the system into web-based or enterprise HR platforms can enable real-time analytics and user-friendly interaction.

Future enhancements may also include advanced explainable AI techniques for deeper interpretability and visualization, as well as cross-domain validation using multiple datasets from different industries. These improvements can increase the scalability, adaptability, and practical applicability of the system.

Moreover, the system can be extended to incorporate automated decision-support mechanisms that recommend specific retention strategies based on predicted attrition risk. By integrating prescriptive analytics, organizations can move beyond prediction to actionable insights, enabling personalized interventions for employees. This advancement would further strengthen the system's role in strategic HR management and contribute to building intelligent, data-driven organizational ecosystems.

Overall, the proposed hybrid framework provides a strong foundation for intelligent employee attrition prediction and can evolve into a comprehensive decision-support system for modern HR analytics.

Chapter 10

Bibliography

10. BIBLIOGRAPHY

- [1] Sirawit Saiwaeo et al., “Developing a Hybrid Machine Learning Model for Employee Turnover Prediction Using LightGBM + GA,” 2023.
- [2] Ahmed, “Predicting Employee Attrition Using Artificial Neural Networks,” 2023.
- [3] J. Judrus et al., “Machine Learning-Based Solution for Predicting Voluntary Employee Turnover,” 2021.
- [4] Z. Zhao, A. Anand, and M. Wang, “Maximum Relevance and Minimum Redundancy Feature Selection Methods,” *Data Science Journal*, vol. 18, no. 1, 2019.
- [5] Z. C. Lipton, “The Mythos of Model Interpretability,” *Communications of the ACM*, vol. 61, no. 10, pp. 36–43, 2018.
- [6] L. Prokhorenkova et al., “CatBoost: Unbiased boosting with categorical features,” *NeurIPS*, 2018.
- [7] Jain and Nayyar, “Predicting Employee Attrition using XGBoost Machine Learning Approach,” 2018.
- [8] G. Ke et al., “LightGBM: A Highly Efficient Gradient Boosting Decision Tree,” *NeurIPS*, 2017.
- [9] S. M. Lundberg and S.-I. Lee, “A Unified Approach to Interpreting Model Predictions,” *NeurIPS*, 2017.
- [10] T. Chen and C. Guestrin, “XGBoost: A Scalable Tree Boosting System,” in *Proc. ACM SIGKDD*, 2016.
- [11] T. Davenport, J. Harris, and J. Shapiro, *Competing on Talent Analytics*. Harvard Business Review Press, 2010.
- [12] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*, 2nd ed. Springer, 2009.

- [13] T. Fawcett, “An Introduction to ROC Analysis,” *Pattern Recognition Letters*, vol. 27, no. 8, pp. 861–874, 2006.
- [14] I. Guyon and A. Elisseeff, “An Introduction to Variable and Feature Selection,” *Journal of Machine Learning Research*, vol. 3, pp. 1157–1182, 2003.
- [15] N. V. Chawla et al., “SMOTE: Synthetic Minority Over-sampling Technique,” *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.
- [16] J. H. Friedman, “Greedy Function Approximation: A Gradient Boosting Machine,” *Annals of Statistics*, vol. 29, no. 5, pp. 1189–1232, 2001.
- [17] L. Breiman, “Random Forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [18] D. E. Goldberg, *Genetic Algorithms in Search, Optimization and Machine Learning*. Addison-Wesley, 1989.